

# project

December 16, 2023

```
[ ]: #Imports

import os
import cv2
import glob
import numpy as np
import random
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader
from xml.etree import ElementTree as et
import albumentations as A
from albumentations.pytorch import ToTensorV2
import matplotlib.pyplot as plt
from torchvision.ops import nms, box_iou
```

```
[ ]: #config

img_size = 512

batch_size = 8

NUM_BOXES = 5
NUM_CLASSES = 5

classes = ['student', 'Security', 'Staff', 'Facility Worker', 'Food Service_
↪worker']

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

print(device)
```

cuda

```
[ ]: #Dataset and DataLoader

class FaceDataset(Dataset):
```

```

def __init__(self, dir_path, width, height, classes, transforms=None):
    self.transforms = transforms
    self.dir_path = dir_path
    self.height = height
    self.num_boxes = 5
    self.num_classes = 5
    self.width = width
    self.classes = classes

    # get all the image paths in sorted order
    self.image_paths = glob.glob(f"{self.dir_path}/*.jpg")
    self.image_paths += glob.glob(f"{self.dir_path}/*.png")
    self.all_images = [image_path.split(
        '/')[-1] for image_path in self.image_paths]
    self.all_images = sorted(self.all_images)

def __getitem__(self, idx):
    # capture the image name and the full image path
    image_name = self.all_images[idx]
    image_path = os.path.join(self.dir_path, image_name)

    # read the image
    image = cv2.imread(image_path)
    # convert BGR to RGB color format
    image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB).astype(np.float32)
    image_resized = cv2.resize(image, (self.width, self.height))
    image_resized /= 255.0

    # capture the corresponding XML file for getting the annotations
    annot_filename = image_name[:-4] + '.xml'
    annot_file_path = os.path.join(self.dir_path, annot_filename)

    boxes = []
    labels = []
    tree = et.parse(annot_file_path)
    root = tree.getroot()

    # get the height and width of the image
    image_width = image.shape[1]
    image_height = image.shape[0]

    # box coordinates for xml files are extracted and corrected for image
    ↪size given
    for member in root.findall('object'):
        # map the current object name to `classes` list to get...
        # ... the label index and append to `labels` list
        labels.append(self.classes.index(member.find('name').text))

```

```

        # xmin = left corner x-coordinates
        xmin = int(member.find('bndbox').find('xmin').text)
        # xmax = right corner x-coordinates
        xmax = int(member.find('bndbox').find('xmax').text)
        # ymin = left corner y-coordinates
        ymin = int(member.find('bndbox').find('ymin').text)
        # ymax = right corner y-coordinates
        ymax = int(member.find('bndbox').find('ymax').text)

        # resize the bounding boxes according to the...
        # ... desired `width`, `height`
        xmin_final = (xmin/image_width)*self.width
        xmax_final = (xmax/image_width)*self.width
        ymin_final = (ymin/image_height)*self.height
        ymax_final = (ymax/image_height)*self.height

        boxes.append([xmin_final, ymin_final, xmax_final, ymax_final])

    boxes, labels = self.map_to_model_output(boxes, labels , self.num_boxes)

    # bounding box to tensor
    boxes = torch.as_tensor(boxes, dtype=torch.float32)
    # labels to tensor
    labels = torch.as_tensor(labels, dtype=torch.int64)

    # prepare the final `target` dictionary
    target = {}
    target["bounding_box"] = boxes
    target["class_label"] = labels

    # apply the image transforms
    if self.transforms:
        sample = self.transforms(image=image_resized,
                                bboxes=target['bounding_box'],
                                labels=labels)

        image_resized = sample['image']
        target['bounding_box'] = torch.Tensor(sample['bboxes'])

    return image_resized, target

def map_to_model_output(self, boxes, classes, num_boxes):
    # Pad the lists to have a fixed number of boxes
    boxes_padded = self.pad_list_to_length(boxes, max_len=num_boxes,
    ↪ pad_value=[0, 0, 512, 512])
    classes_padded = self.pad_list_to_length(classes, max_len=num_boxes,
    ↪ pad_value=-1)

```

```

        return boxes_padded, classes_padded

    def pad_list_to_length(self, list_to_pad, max_len, pad_value):
        list_len = len(list_to_pad)
        if list_len >= max_len:
            return list_to_pad[:max_len]
        else:
            # Calculate padding required
            padding = [pad_value] * (max_len - list_len)

            # Pad the list
            padded_list = list_to_pad + padding

            return padded_list

    def __len__(self):
        return len(self.all_images)

train_transforms = A.Compose([
    A.Flip(0.5),
    A.RandomRotate90(0.5),
    A.MotionBlur(p=0.2),
    A.MedianBlur(blur_limit=3, p=0.1),
    A.Blur(blur_limit=3, p=0.1),
    ToTensorV2(p=1.0),
], bbox_params={
    'format': 'pascal_voc',
    'label_fields': ['labels']
})

valid_transforms = A.Compose([ToTensorV2(p=1.0)], bbox_params={'format': 'pascal_voc', 'label_fields': ['labels']})

def collate_fn(batch):
    images, targets = zip(*batch)

    bounding_boxes = [target['bounding_box'] for target in targets]
    class_labels = [target['class_label'] for target in targets]

    return images, bounding_boxes, class_labels

train_dir = 'dataset/train'
valid_dir = 'dataset/valid'

train_dataset = FaceDataset(train_dir, img_size, img_size, classes, transforms=train_transforms)

```

```

valid_dataset =
    ↳FaceDataset(valid_dir,img_size,img_size,classes,transforms=valid_transforms)

train_loader = DataLoader(
    train_dataset,
    batch_size=batch_size,
    shuffle=True,
    num_workers=0,
    collate_fn=collate_fn
)

valid_loader = DataLoader(
    valid_dataset,
    batch_size=batch_size,
    shuffle=False,
    num_workers=0,
    collate_fn=collate_fn
)

print("Input Size")
for images, bounding_boxes, class_labels in train_loader:
    print(images[0].shape)
    print(bounding_boxes[0].shape)
    print(class_labels[0].shape)
    break

```

```

Input Size
torch.Size([3, 512, 512])
torch.Size([5, 4])
torch.Size([5])

```

[ ]: *#Model Definition*

```

class ResNet18(nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1 = nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3)
        self.bn1 = nn.BatchNorm2d(64)
        self.relu = nn.ReLU(inplace=True)

        # Max pooling for 512x512 images
        self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)

```

```

        # Adjust feature map sizes for 512 input
        self.res1 = ResNetBlock(64, 128, 128)
        self.res2 = ResNetBlock(128, 256, 256)
        self.res3 = ResNetBlock(256, 256, 512)
        self.res4 = ResNetBlock(256, 512, 512)

        self.avgpool = nn.AdaptiveAvgPool2d((1, 1))

    def forward(self, x):
        x = self.conv1(x)
        x = self.bn1(x)
        x = self.relu(x)

        # Max pooling
        x = self.maxpool(x)

        x = self.res1(x)
        x = self.res2(x)
        x = self.res3(x)
        x = self.res4(x)

        x = self.avgpool(x)

        return x

class ResNetBlock(nn.Module):
    def __init__(self, in_channels, out_channels, hidden_dim):
        super().__init__()

        self.layers = nn.Sequential(
            nn.Conv2d(in_channels, hidden_dim, kernel_size=3, padding=1),
            nn.BatchNorm2d(hidden_dim),
            nn.ReLU(inplace=True),
            nn.Conv2d(hidden_dim, out_channels, kernel_size=3, padding=1),
            nn.BatchNorm2d(out_channels)
        )

    def forward(self, x):
        # Fixed: Identity mapping
        identity = x
        out = self.layers(x)

        # Ensure identity has the same dimensions as out
        if identity.size() != out.size():
            identity = F.pad(identity, (0, 0, 0, 0, 0, 0, out.size(1) - identity.
↪size(1)))

```

```

        out += identity
        out = nn.ReLU(inplace=True)(out)
        return out

class FaceRecogNet(nn.Module):

    def __init__(self):
        super().__init__()

        self.cnn = ResNet18()

        self.bbox_reg = nn.Linear(512, NUM_BOXES*4)

        self.cls_head = nn.Linear(512, NUM_CLASSES)

    def forward(self, x):

        x = self.cnn(x)

        x = x.view(x.size(0), -1)

        predicted_boxes = self.bbox_reg(x)
        predicted_classes = self.cls_head(x)

        predicted_boxes = predicted_boxes.reshape(x.size(0), NUM_BOXES, 4)
        predicted_classes = predicted_classes.reshape(x.size(0), NUM_CLASSES)

        return predicted_boxes, predicted_classes

```

[ ]: *#Objective Function*

```

def Calculate_loss(anchor, positive, negative):

    # Calculate the loss for the bounding boxes
    loss_bbox = F.triplet_margin_loss(anchor[0], positive[0], negative[0],
    ↪margin = 0.5, p = 4, reduction='mean')

    loss_class = F.triplet_margin_with_distance_loss(anchor[0], positive[0],
    ↪negative[0], distance_function=nn.PairwiseDistance(), margin= 0.1
    ↪, reduction='mean')

    total_loss = loss_bbox + loss_class

    return total_loss

```

```
[ ]: def train_tune(config, out_dir='outputs'):

    model = FaceRecogNet().to(device)

    def init_weights(m):
        if isinstance(m, nn.Conv2d):
            torch.nn.init.kaiming_normal_(m.weight)
            torch.nn.init.zeros_(m.bias)
        if isinstance(m, nn.Linear):
            torch.nn.init.kaiming_normal_(m.weight)
            torch.nn.init.zeros_(m.bias)

    model.apply(init_weights)

    optimizer = torch.optim.SGD(model.parameters(), lr=config['lr'], momentum =
    ↪config['momentum'], weight_decay = 0.0005)

    lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=2,
    ↪gamma=0.1, last_epoch=-1)

    train_losses = []
    valid_losses = []

    num_epochs = config["epochs"]
    for epoch in range(num_epochs):
        model.train()
        for images, bounding_boxes, class_labels in train_loader:
            images = torch.stack(images).to(device)
            bounding_boxes = torch.stack(bounding_boxes).to(device)
            class_labels = torch.stack(class_labels).to(device)

            optimizer.zero_grad()

            anchor = model(images)

            positive = model(images)

            negative = model(images)

            # Calculate loss
            loss = Calculate_loss(anchor, positive, negative)

            # Backpropagation
            loss.backward()

            # Update weights
            optimizer.step()
```



```

        train_losses.append(loss.item())

    print(f"Epoch: {epoch+1}/{num_epochs}, Training Loss: {loss.item():.
↪4f}")

    # Validation
    model.eval()
    with torch.no_grad():
        for images, bounding_boxes, class_labels in valid_loader:

            images = torch.stack(images).to(device)
            bounding_boxes = torch.stack(bounding_boxes).to(device)
            class_labels = torch.stack(class_labels).to(device)

            anchor_valid = model(images)
            positive_valid = model(images)
            negative_valid = model(images)

            # Calculate loss
            loss = _
↪Calculate_loss(anchor_valid, positive_valid, negative_valid)

            valid_losses.append(loss.item())

    print(f"Epoch: {epoch+1}/{num_epochs}, Validation Loss: {loss.item():.
↪4f}")

    lr_scheduler.step()

plt.plot(train_losses, label='Training loss')
plt.plot(valid_losses, label='Validation loss')
plt.legend(frameon=False)
plt.savefig( f"{out_dir}/_loss_plot_{num_epochs}.png")
torch.save(model.state_dict(), f"{out_dir}/_{num_epochs}.pth")

```

[ ]: *#Hyperparameter Tuning*

```

hyperparameter_space = {
    "lr": [1e-7, 2e-9],
    "momentum": [0.7, 0.2],
    "epochs": [10, 15, 20]
}

```

```

for lr in hyperparameter_space["lr"]:
    for momentum in hyperparameter_space["momentum"]:
        for epochs in hyperparameter_space["epochs"]:
            config = {"lr": lr, "momentum": momentum, "epochs": epochs}
            print("Training with config: ", config)
            train_tune(config, out_dir='outputs')
            torch.cuda.empty_cache()

```

Training with config: {'lr': 1e-07, 'momentum': 0.7, 'epochs': 10}

Epoch: 1/10, Training Loss: 0.6000  
 Epoch: 1/10, Validation Loss: 0.6000  
 Epoch: 2/10, Training Loss: 0.6000  
 Epoch: 2/10, Validation Loss: 0.6000  
 Epoch: 3/10, Training Loss: 0.6000  
 Epoch: 3/10, Validation Loss: 0.6000  
 Epoch: 4/10, Training Loss: 0.6000  
 Epoch: 4/10, Validation Loss: 0.6000  
 Epoch: 5/10, Training Loss: 0.6000  
 Epoch: 5/10, Validation Loss: 0.6000  
 Epoch: 6/10, Training Loss: 0.6000  
 Epoch: 6/10, Validation Loss: 0.6000  
 Epoch: 7/10, Training Loss: 0.6000  
 Epoch: 7/10, Validation Loss: 0.6000  
 Epoch: 8/10, Training Loss: 0.6000  
 Epoch: 8/10, Validation Loss: 0.6000  
 Epoch: 9/10, Training Loss: 0.6000  
 Epoch: 9/10, Validation Loss: 0.6000  
 Epoch: 10/10, Training Loss: 0.6000  
 Epoch: 10/10, Validation Loss: 0.6000

Training with config: {'lr': 1e-07, 'momentum': 0.7, 'epochs': 15}

Epoch: 1/15, Training Loss: 0.6000  
 Epoch: 1/15, Validation Loss: 0.6000  
 Epoch: 2/15, Training Loss: 0.6000  
 Epoch: 2/15, Validation Loss: 0.6000  
 Epoch: 3/15, Training Loss: 0.6000  
 Epoch: 3/15, Validation Loss: 0.6000  
 Epoch: 4/15, Training Loss: 0.6000  
 Epoch: 4/15, Validation Loss: 0.6000  
 Epoch: 5/15, Training Loss: 0.6000  
 Epoch: 5/15, Validation Loss: 0.6000  
 Epoch: 6/15, Training Loss: 0.6000  
 Epoch: 6/15, Validation Loss: 0.6000  
 Epoch: 7/15, Training Loss: 0.6000  
 Epoch: 7/15, Validation Loss: 0.6000  
 Epoch: 8/15, Training Loss: 0.6000  
 Epoch: 8/15, Validation Loss: 0.6000

Epoch: 9/15, Training Loss: 0.6000  
Epoch: 9/15, Validation Loss: 0.6000  
Epoch: 10/15, Training Loss: 0.6000  
Epoch: 10/15, Validation Loss: 0.6000  
Epoch: 11/15, Training Loss: 0.6000  
Epoch: 11/15, Validation Loss: 0.6000  
Epoch: 12/15, Training Loss: 0.6000  
Epoch: 12/15, Validation Loss: 0.6000  
Epoch: 13/15, Training Loss: 0.6000  
Epoch: 13/15, Validation Loss: 0.6000  
Epoch: 14/15, Training Loss: 0.6000  
Epoch: 14/15, Validation Loss: 0.6000  
Epoch: 15/15, Training Loss: 0.6000  
Epoch: 15/15, Validation Loss: 0.6000  
Training with config: {'lr': 1e-07, 'momentum': 0.7, 'epochs': 20}  
Epoch: 1/20, Training Loss: 0.6000  
Epoch: 1/20, Validation Loss: 0.6000  
Epoch: 2/20, Training Loss: 0.6000  
Epoch: 2/20, Validation Loss: 0.6000  
Epoch: 3/20, Training Loss: 0.6000  
Epoch: 3/20, Validation Loss: 0.6000  
Epoch: 4/20, Training Loss: 0.6000  
Epoch: 4/20, Validation Loss: 0.6000  
Epoch: 5/20, Training Loss: 0.6000  
Epoch: 5/20, Validation Loss: 0.6000  
Epoch: 6/20, Training Loss: 0.6000  
Epoch: 6/20, Validation Loss: 0.6000  
Epoch: 7/20, Training Loss: 0.6000  
Epoch: 7/20, Validation Loss: 0.6000  
Epoch: 8/20, Training Loss: 0.6000  
Epoch: 8/20, Validation Loss: 0.6000  
Epoch: 9/20, Training Loss: 0.6000  
Epoch: 9/20, Validation Loss: 0.6000  
Epoch: 10/20, Training Loss: 0.6000  
Epoch: 10/20, Validation Loss: 0.6000  
Epoch: 11/20, Training Loss: 0.6000  
Epoch: 11/20, Validation Loss: 0.6000  
Epoch: 12/20, Training Loss: 0.6000  
Epoch: 12/20, Validation Loss: 0.6000  
Epoch: 13/20, Training Loss: 0.6000  
Epoch: 13/20, Validation Loss: 0.6000  
Epoch: 14/20, Training Loss: 0.6000  
Epoch: 14/20, Validation Loss: 0.6000  
Epoch: 15/20, Training Loss: 0.6000  
Epoch: 15/20, Validation Loss: 0.6000  
Epoch: 16/20, Training Loss: 0.6000  
Epoch: 16/20, Validation Loss: 0.6000  
Epoch: 17/20, Training Loss: 0.6000

Epoch: 17/20, Validation Loss: 0.6000  
Epoch: 18/20, Training Loss: 0.6000  
Epoch: 18/20, Validation Loss: 0.6000  
Epoch: 19/20, Training Loss: 0.6000  
Epoch: 19/20, Validation Loss: 0.6000  
Epoch: 20/20, Training Loss: 0.6000  
Epoch: 20/20, Validation Loss: 0.6000  
Training with config: {'lr': 1e-07, 'momentum': 0.2, 'epochs': 10}  
Epoch: 1/10, Training Loss: 0.6000  
Epoch: 1/10, Validation Loss: 0.6000  
Epoch: 2/10, Training Loss: 0.6000  
Epoch: 2/10, Validation Loss: 0.6000  
Epoch: 3/10, Training Loss: 0.6000  
Epoch: 3/10, Validation Loss: 0.6000  
Epoch: 4/10, Training Loss: 0.6000  
Epoch: 4/10, Validation Loss: 0.6000  
Epoch: 5/10, Training Loss: 0.6000  
Epoch: 5/10, Validation Loss: 0.6000  
Epoch: 6/10, Training Loss: 0.6000  
Epoch: 6/10, Validation Loss: 0.6000  
Epoch: 7/10, Training Loss: 0.6000  
Epoch: 7/10, Validation Loss: 0.6000  
Epoch: 8/10, Training Loss: 0.6000  
Epoch: 8/10, Validation Loss: 0.6000  
Epoch: 9/10, Training Loss: 0.6000  
Epoch: 9/10, Validation Loss: 0.6000  
Epoch: 10/10, Training Loss: 0.6000  
Epoch: 10/10, Validation Loss: 0.6000  
Training with config: {'lr': 1e-07, 'momentum': 0.2, 'epochs': 15}  
Epoch: 1/15, Training Loss: 0.6000  
Epoch: 1/15, Validation Loss: 0.6000  
Epoch: 2/15, Training Loss: 0.6000  
Epoch: 2/15, Validation Loss: 0.6000  
Epoch: 3/15, Training Loss: 0.6000  
Epoch: 3/15, Validation Loss: 0.6000  
Epoch: 4/15, Training Loss: 0.6000  
Epoch: 4/15, Validation Loss: 0.6000  
Epoch: 5/15, Training Loss: 0.6000  
Epoch: 5/15, Validation Loss: 0.6000  
Epoch: 6/15, Training Loss: 0.6000  
Epoch: 6/15, Validation Loss: 0.6000  
Epoch: 7/15, Training Loss: 0.6000  
Epoch: 7/15, Validation Loss: 0.6000  
Epoch: 8/15, Training Loss: 0.6000  
Epoch: 8/15, Validation Loss: 0.6000  
Epoch: 9/15, Training Loss: 0.6000  
Epoch: 9/15, Validation Loss: 0.6000  
Epoch: 10/15, Training Loss: 0.6000

Epoch: 10/15, Validation Loss: 0.6000  
Epoch: 11/15, Training Loss: 0.6000  
Epoch: 11/15, Validation Loss: 0.6000  
Epoch: 12/15, Training Loss: 0.6000  
Epoch: 12/15, Validation Loss: 0.6000  
Epoch: 13/15, Training Loss: 0.6000  
Epoch: 13/15, Validation Loss: 0.6000  
Epoch: 14/15, Training Loss: 0.6000  
Epoch: 14/15, Validation Loss: 0.6000  
Epoch: 15/15, Training Loss: 0.6000  
Epoch: 15/15, Validation Loss: 0.6000  
Training with config: {'lr': 1e-07, 'momentum': 0.2, 'epochs': 20}  
Epoch: 1/20, Training Loss: 0.6000  
Epoch: 1/20, Validation Loss: 0.6000  
Epoch: 2/20, Training Loss: 0.6000  
Epoch: 2/20, Validation Loss: 0.6000  
Epoch: 3/20, Training Loss: 0.6000  
Epoch: 3/20, Validation Loss: 0.6000  
Epoch: 4/20, Training Loss: 0.6000  
Epoch: 4/20, Validation Loss: 0.6000  
Epoch: 5/20, Training Loss: 0.6000  
Epoch: 5/20, Validation Loss: 0.6000  
Epoch: 6/20, Training Loss: 0.6000  
Epoch: 6/20, Validation Loss: 0.6000  
Epoch: 7/20, Training Loss: 0.6000  
Epoch: 7/20, Validation Loss: 0.6000  
Epoch: 8/20, Training Loss: 0.6000  
Epoch: 8/20, Validation Loss: 0.6000  
Epoch: 9/20, Training Loss: 0.6000  
Epoch: 9/20, Validation Loss: 0.6000  
Epoch: 10/20, Training Loss: 0.6000  
Epoch: 10/20, Validation Loss: 0.6000  
Epoch: 11/20, Training Loss: 0.6000  
Epoch: 11/20, Validation Loss: 0.6000  
Epoch: 12/20, Training Loss: 0.6000  
Epoch: 12/20, Validation Loss: 0.6000  
Epoch: 13/20, Training Loss: 0.6000  
Epoch: 13/20, Validation Loss: 0.6000  
Epoch: 14/20, Training Loss: 0.6000  
Epoch: 14/20, Validation Loss: 0.6000  
Epoch: 15/20, Training Loss: 0.6000  
Epoch: 15/20, Validation Loss: 0.6000  
Epoch: 16/20, Training Loss: 0.6000  
Epoch: 16/20, Validation Loss: 0.6000  
Epoch: 17/20, Training Loss: 0.6000  
Epoch: 17/20, Validation Loss: 0.6000  
Epoch: 18/20, Training Loss: 0.6000  
Epoch: 18/20, Validation Loss: 0.6000

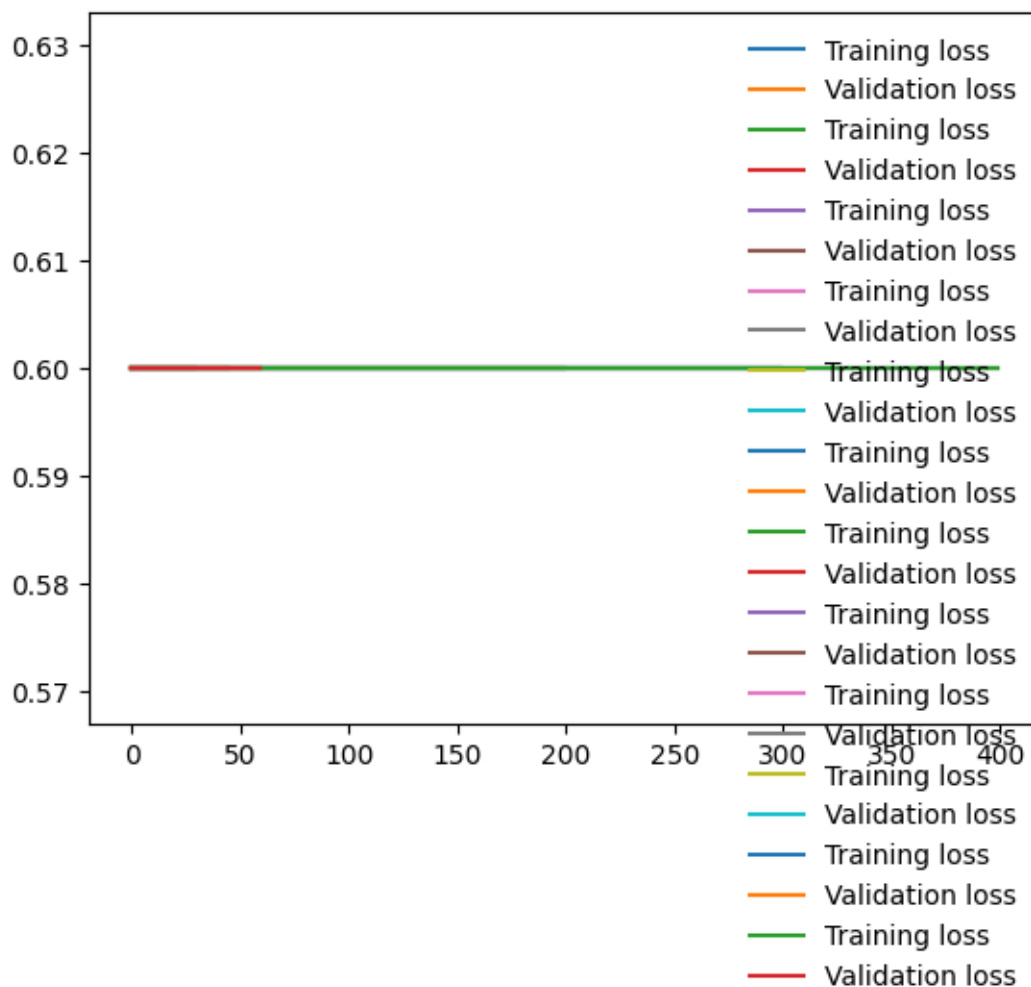
Epoch: 19/20, Training Loss: 0.6000  
Epoch: 19/20, Validation Loss: 0.6000  
Epoch: 20/20, Training Loss: 0.6000  
Epoch: 20/20, Validation Loss: 0.6000  
Training with config: {'lr': 2e-09, 'momentum': 0.7, 'epochs': 10}  
Epoch: 1/10, Training Loss: 0.6000  
Epoch: 1/10, Validation Loss: 0.6000  
Epoch: 2/10, Training Loss: 0.6000  
Epoch: 2/10, Validation Loss: 0.6000  
Epoch: 3/10, Training Loss: 0.6000  
Epoch: 3/10, Validation Loss: 0.6000  
Epoch: 4/10, Training Loss: 0.6000  
Epoch: 4/10, Validation Loss: 0.6000  
Epoch: 5/10, Training Loss: 0.6000  
Epoch: 5/10, Validation Loss: 0.6000  
Epoch: 6/10, Training Loss: 0.6000  
Epoch: 6/10, Validation Loss: 0.6000  
Epoch: 7/10, Training Loss: 0.6000  
Epoch: 7/10, Validation Loss: 0.6000  
Epoch: 8/10, Training Loss: 0.6000  
Epoch: 8/10, Validation Loss: 0.6000  
Epoch: 9/10, Training Loss: 0.6000  
Epoch: 9/10, Validation Loss: 0.6000  
Epoch: 10/10, Training Loss: 0.6000  
Epoch: 10/10, Validation Loss: 0.6000  
Training with config: {'lr': 2e-09, 'momentum': 0.7, 'epochs': 15}  
Epoch: 1/15, Training Loss: 0.6000  
Epoch: 1/15, Validation Loss: 0.6000  
Epoch: 2/15, Training Loss: 0.6000  
Epoch: 2/15, Validation Loss: 0.6000  
Epoch: 3/15, Training Loss: 0.6000  
Epoch: 3/15, Validation Loss: 0.6000  
Epoch: 4/15, Training Loss: 0.6000  
Epoch: 4/15, Validation Loss: 0.6000  
Epoch: 5/15, Training Loss: 0.6000  
Epoch: 5/15, Validation Loss: 0.6000  
Epoch: 6/15, Training Loss: 0.6000  
Epoch: 6/15, Validation Loss: 0.6000  
Epoch: 7/15, Training Loss: 0.6000  
Epoch: 7/15, Validation Loss: 0.6000  
Epoch: 8/15, Training Loss: 0.6000  
Epoch: 8/15, Validation Loss: 0.6000  
Epoch: 9/15, Training Loss: 0.6000  
Epoch: 9/15, Validation Loss: 0.6000  
Epoch: 10/15, Training Loss: 0.6000  
Epoch: 10/15, Validation Loss: 0.6000  
Epoch: 11/15, Training Loss: 0.6000  
Epoch: 11/15, Validation Loss: 0.6000

Epoch: 12/15, Training Loss: 0.6000  
Epoch: 12/15, Validation Loss: 0.6000  
Epoch: 13/15, Training Loss: 0.6000  
Epoch: 13/15, Validation Loss: 0.6000  
Epoch: 14/15, Training Loss: 0.6000  
Epoch: 14/15, Validation Loss: 0.6000  
Epoch: 15/15, Training Loss: 0.6000  
Epoch: 15/15, Validation Loss: 0.6000  
Training with config: {'lr': 2e-09, 'momentum': 0.7, 'epochs': 20}  
Epoch: 1/20, Training Loss: 0.6000  
Epoch: 1/20, Validation Loss: 0.6000  
Epoch: 2/20, Training Loss: 0.6000  
Epoch: 2/20, Validation Loss: 0.6000  
Epoch: 3/20, Training Loss: 0.6000  
Epoch: 3/20, Validation Loss: 0.6000  
Epoch: 4/20, Training Loss: 0.6000  
Epoch: 4/20, Validation Loss: 0.6000  
Epoch: 5/20, Training Loss: 0.6000  
Epoch: 5/20, Validation Loss: 0.6000  
Epoch: 6/20, Training Loss: 0.6000  
Epoch: 6/20, Validation Loss: 0.6000  
Epoch: 7/20, Training Loss: 0.6000  
Epoch: 7/20, Validation Loss: 0.6000  
Epoch: 8/20, Training Loss: 0.6000  
Epoch: 8/20, Validation Loss: 0.6000  
Epoch: 9/20, Training Loss: 0.6000  
Epoch: 9/20, Validation Loss: 0.6000  
Epoch: 10/20, Training Loss: 0.6000  
Epoch: 10/20, Validation Loss: 0.6000  
Epoch: 11/20, Training Loss: 0.6000  
Epoch: 11/20, Validation Loss: 0.6000  
Epoch: 12/20, Training Loss: 0.6000  
Epoch: 12/20, Validation Loss: 0.6000  
Epoch: 13/20, Training Loss: 0.6000  
Epoch: 13/20, Validation Loss: 0.6000  
Epoch: 14/20, Training Loss: 0.6000  
Epoch: 14/20, Validation Loss: 0.6000  
Epoch: 15/20, Training Loss: 0.6000  
Epoch: 15/20, Validation Loss: 0.6000  
Epoch: 16/20, Training Loss: 0.6000  
Epoch: 16/20, Validation Loss: 0.6000  
Epoch: 17/20, Training Loss: 0.6000  
Epoch: 17/20, Validation Loss: 0.6000  
Epoch: 18/20, Training Loss: 0.6000  
Epoch: 18/20, Validation Loss: 0.6000  
Epoch: 19/20, Training Loss: 0.6000  
Epoch: 19/20, Validation Loss: 0.6000  
Epoch: 20/20, Training Loss: 0.6000

Epoch: 20/20, Validation Loss: 0.6000  
Training with config: {'lr': 2e-09, 'momentum': 0.2, 'epochs': 10}  
Epoch: 1/10, Training Loss: 0.6000  
Epoch: 1/10, Validation Loss: 0.6000  
Epoch: 2/10, Training Loss: 0.6000  
Epoch: 2/10, Validation Loss: 0.6000  
Epoch: 3/10, Training Loss: 0.6000  
Epoch: 3/10, Validation Loss: 0.6000  
Epoch: 4/10, Training Loss: 0.6000  
Epoch: 4/10, Validation Loss: 0.6000  
Epoch: 5/10, Training Loss: 0.6000  
Epoch: 5/10, Validation Loss: 0.6000  
Epoch: 6/10, Training Loss: 0.6000  
Epoch: 6/10, Validation Loss: 0.6000  
Epoch: 7/10, Training Loss: 0.6000  
Epoch: 7/10, Validation Loss: 0.6000  
Epoch: 8/10, Training Loss: 0.6000  
Epoch: 8/10, Validation Loss: 0.6000  
Epoch: 9/10, Training Loss: 0.6000  
Epoch: 9/10, Validation Loss: 0.6000  
Epoch: 10/10, Training Loss: 0.6000  
Epoch: 10/10, Validation Loss: 0.6000  
Training with config: {'lr': 2e-09, 'momentum': 0.2, 'epochs': 15}  
Epoch: 1/15, Training Loss: 0.6000  
Epoch: 1/15, Validation Loss: 0.6000  
Epoch: 2/15, Training Loss: 0.6000  
Epoch: 2/15, Validation Loss: 0.6000  
Epoch: 3/15, Training Loss: 0.6000  
Epoch: 3/15, Validation Loss: 0.6000  
Epoch: 4/15, Training Loss: 0.6000  
Epoch: 4/15, Validation Loss: 0.6000  
Epoch: 5/15, Training Loss: 0.6000  
Epoch: 5/15, Validation Loss: 0.6000  
Epoch: 6/15, Training Loss: 0.6000  
Epoch: 6/15, Validation Loss: 0.6000  
Epoch: 7/15, Training Loss: 0.6000  
Epoch: 7/15, Validation Loss: 0.6000  
Epoch: 8/15, Training Loss: 0.6000  
Epoch: 8/15, Validation Loss: 0.6000  
Epoch: 9/15, Training Loss: 0.6000  
Epoch: 9/15, Validation Loss: 0.6000  
Epoch: 10/15, Training Loss: 0.6000  
Epoch: 10/15, Validation Loss: 0.6000  
Epoch: 11/15, Training Loss: 0.6000  
Epoch: 11/15, Validation Loss: 0.6000  
Epoch: 12/15, Training Loss: 0.6000  
Epoch: 12/15, Validation Loss: 0.6000  
Epoch: 13/15, Training Loss: 0.6000



Epoch: 13/15, Validation Loss: 0.6000  
Epoch: 14/15, Training Loss: 0.6000  
Epoch: 14/15, Validation Loss: 0.6000  
Epoch: 15/15, Training Loss: 0.6000  
Epoch: 15/15, Validation Loss: 0.6000  
Training with config: {'lr': 2e-09, 'momentum': 0.2, 'epochs': 20}  
Epoch: 1/20, Training Loss: 0.6000  
Epoch: 1/20, Validation Loss: 0.6000  
Epoch: 2/20, Training Loss: 0.6000  
Epoch: 2/20, Validation Loss: 0.6000  
Epoch: 3/20, Training Loss: 0.6000  
Epoch: 3/20, Validation Loss: 0.6000  
Epoch: 4/20, Training Loss: 0.6000  
Epoch: 4/20, Validation Loss: 0.6000  
Epoch: 5/20, Training Loss: 0.6000  
Epoch: 5/20, Validation Loss: 0.6000  
Epoch: 6/20, Training Loss: 0.6000  
Epoch: 6/20, Validation Loss: 0.6000  
Epoch: 7/20, Training Loss: 0.6000  
Epoch: 7/20, Validation Loss: 0.6000  
Epoch: 8/20, Training Loss: 0.6000  
Epoch: 8/20, Validation Loss: 0.6000  
Epoch: 9/20, Training Loss: 0.6000  
Epoch: 9/20, Validation Loss: 0.6000  
Epoch: 10/20, Training Loss: 0.6000  
Epoch: 10/20, Validation Loss: 0.6000  
Epoch: 11/20, Training Loss: 0.6000  
Epoch: 11/20, Validation Loss: 0.6000  
Epoch: 12/20, Training Loss: 0.6000  
Epoch: 12/20, Validation Loss: 0.6000  
Epoch: 13/20, Training Loss: 0.6000  
Epoch: 13/20, Validation Loss: 0.6000  
Epoch: 14/20, Training Loss: 0.6000  
Epoch: 14/20, Validation Loss: 0.6000  
Epoch: 15/20, Training Loss: 0.6000  
Epoch: 15/20, Validation Loss: 0.6000  
Epoch: 16/20, Training Loss: 0.6000  
Epoch: 16/20, Validation Loss: 0.6000  
Epoch: 17/20, Training Loss: 0.6000  
Epoch: 17/20, Validation Loss: 0.6000  
Epoch: 18/20, Training Loss: 0.6000  
Epoch: 18/20, Validation Loss: 0.6000  
Epoch: 19/20, Training Loss: 0.6000  
Epoch: 19/20, Validation Loss: 0.6000  
Epoch: 20/20, Training Loss: 0.6000  
Epoch: 20/20, Validation Loss: 0.6000



```
[ ]: #Testing the Recall and Preicision score for the best model

model = FaceRecogNet().to('cpu')

model.load_state_dict(torch.load('outputs/model.pth',map_location=torch.
    ↪device('cpu'))))

[ ]: <All keys matched successfully>

[ ]: #Calculate the IOU score for the predicted bounding boxes and the ground truth

def calculate_iou(box1, box2):
    intersection = box_iou(box1.unsqueeze(0), box2.unsqueeze(0))
    return intersection.item()

def calculate_precision_recall(targets, predictions, iou_threshold=0.5):
```

```

true_positives = 0
false_positives = 0
false_negatives = 0

for batch_idx in range(len(predictions['boxes'])):
    pred_boxes = predictions['boxes'][batch_idx]
    pred_labels = predictions['labels'][batch_idx]

    for pred_box, pred_label in zip(pred_boxes, pred_labels):
        iou_max = 0
        matching_target_index = None

        for i, (target_box, target_label) in
↪ enumerate(zip(targets['boxes'][batch_idx], targets['labels'][batch_idx])):
            iou = calculate_iou(pred_box, target_box)
            if iou > iou_threshold and iou > iou_max:
                iou_max = iou
                matching_target_index = i

        if matching_target_index is not None:
            true_positives += 1
            targets['boxes'][batch_idx] = torch.
↪ cat([targets['boxes'][batch_idx][:matching_target_index],
↪ targets['boxes'][batch_idx][matching_target_index+1:]])
            targets['labels'][batch_idx] = torch.
↪ cat([targets['labels'][batch_idx][:matching_target_index],
↪ targets['labels'][batch_idx][matching_target_index+1:]])
        else:
            false_positives += 1

    false_negatives = sum(len(targets['boxes'][batch_idx]) for batch_idx in
↪ range(len(targets['boxes'])))

    precision = true_positives / (true_positives + false_positives) if
↪ (true_positives + false_positives) > 0 else 0
    recall = true_positives / (true_positives + false_negatives) if
↪ (true_positives + false_negatives) > 0 else 0

    return precision, recall

```

[ ]: *#Getting the best model and calculating precision and recall based on IOU score*

```

target_images, target_boxes, target_labels = next(iter(valid_loader))

target_images = torch.stack(target_images).to('cpu')
target_boxes = torch.stack(target_boxes).to('cpu')

```

```
target_labels = torch.stack(target_labels).to('cpu')

prediction_boxes, prediction_labels = model(target_images)

predictions = {'boxes': prediction_boxes , 'labels': prediction_labels}

targets = {'boxes': target_boxes , 'labels': target_labels}
precision, recall = calculate_precision_recall(targets, predictions)

print(f'Precision: {precision:.4f}, Recall: {recall:.4f}')
```

Precision: 0.4, Recall: 0.6